Microsoft

Como ensinar uma máquina a escrever com Deep Learning, CNTK, TensorFlow e Azure

2018 Global Azure BOOTCAMP

Vitor Meriat

Nossos agradecimentos aos patrocinadores globais de 2018!

E um GRANDE agradecimento ao nosso patrocinador local!

venturus

inovação & tecnologia

http://www.venturus.org.br

About me

Vitor is a computer scientist who is passionate about creating software that will positively change the world we live in.

Currently, he works as **Technical Evangelist** at **ESX**, where he is helping to shape new disruptive services based in Cloud Computing. Data science enthusiast, he works with Big Data projects, data analytics and **Microsoft MVP Azure**.

vitormeriat.com.br

linkedin.com/in/vitormeriat

twitter.com/vitormeriat

github.com/vitormeriat

youtube.com/vitormeriat

esx

2018 Global Azure BOOTCAMP

Microsoft

Microsoft® MVP Most Valuable Professional

Agenda

O que é Visão Computacional e por que é uma assunto tão difícil;

Como eu ensino um computador a escrever?

Como Deep Learning e Cloud me ajudam com isso?

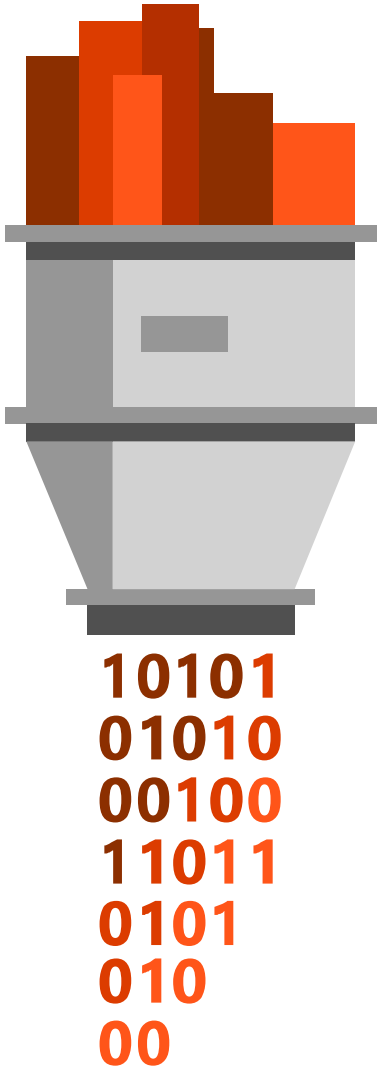

A faster, more efficient, more intelligent cloud

The need for **SCALE**

Data explosion: 2013 4.4 ZB - 2020 44 ZB
ML, DNN, AI are driving requirements up faster

The need for **LOW-LATENCY**

Autonomous decision making
Real-time insights into connected devices
Interactive user experiences

The need for **THROUGHPUT**

Cloud-scale services
Searches and recommendations (Indexing the Internet!)

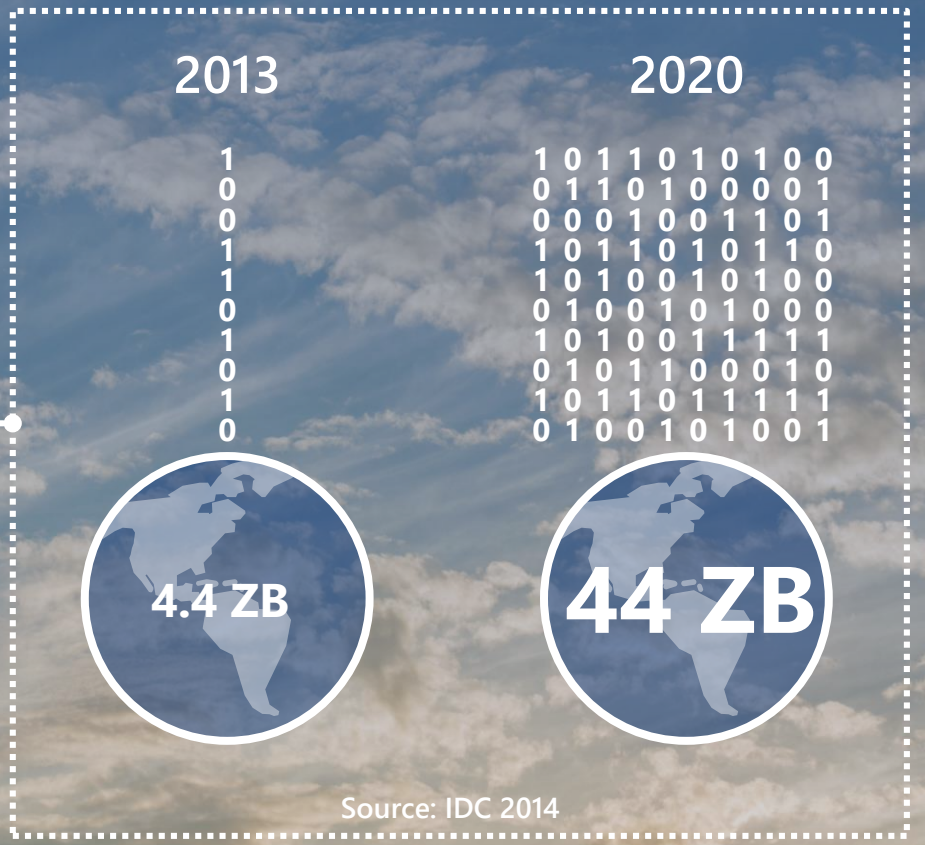

New Azure VM Sizes

Lowest Price

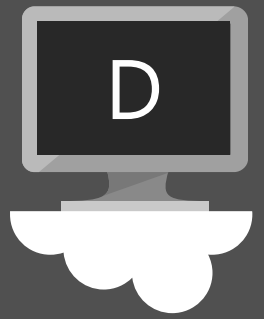

SSD Storage
Fast CPUs

New generation
of D family VMs

High memory and
Large SSDs

New A-Series

Compute Intensive

NVIDIA GPUs
K80 Compute

NVIDIA GPUs
M60 Visualization

Fastest CPU
IB Connectivity

Large SSDs

SAP Large Instances

Azure VM Sizes

Sign up for preview:
https://aka.ms/gpupreviewndnc

Azure Batch AI Training

AI Training At Scale with Azure

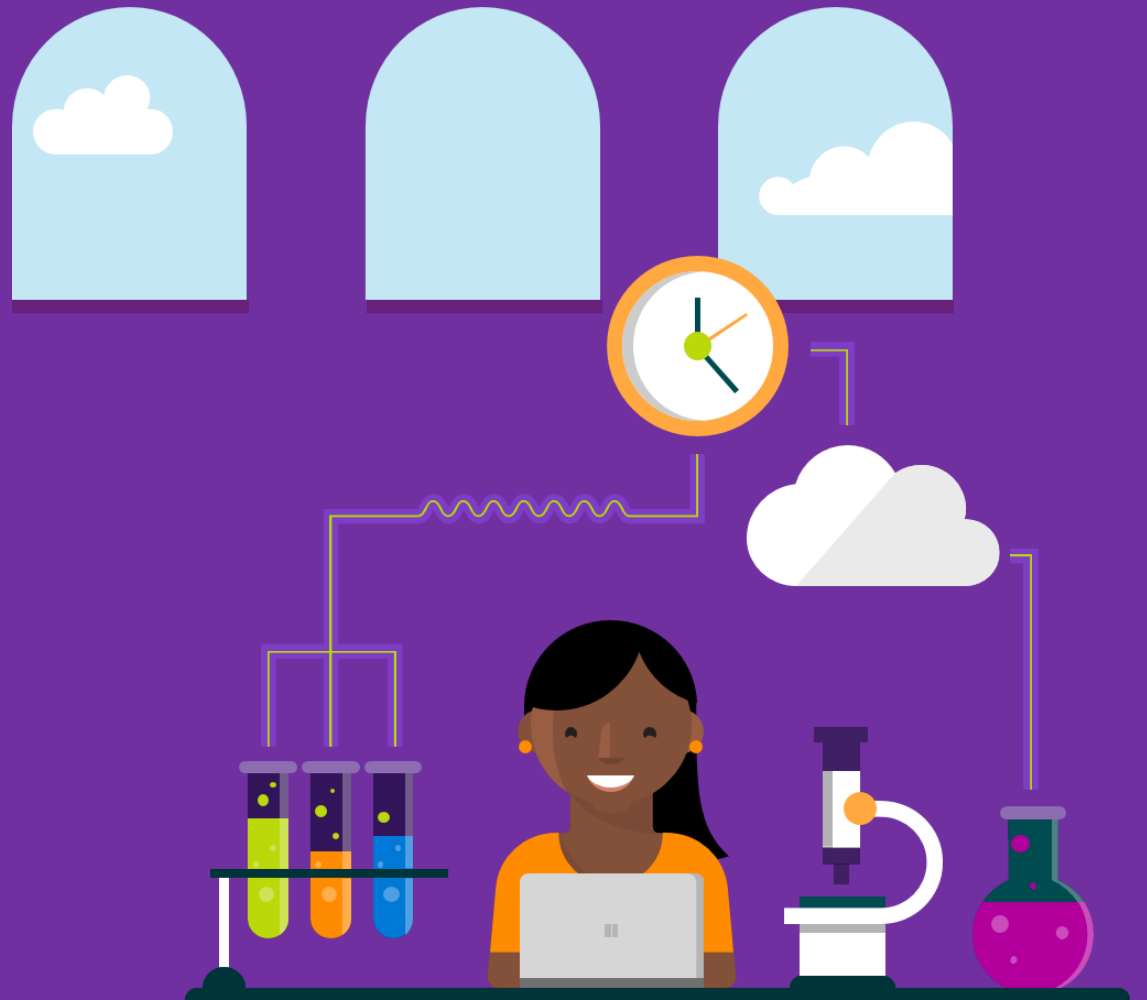

Announcing
Azure Batch Rendering Service

Autodesk 3ds Max / Maya

Integrated Client Plugin

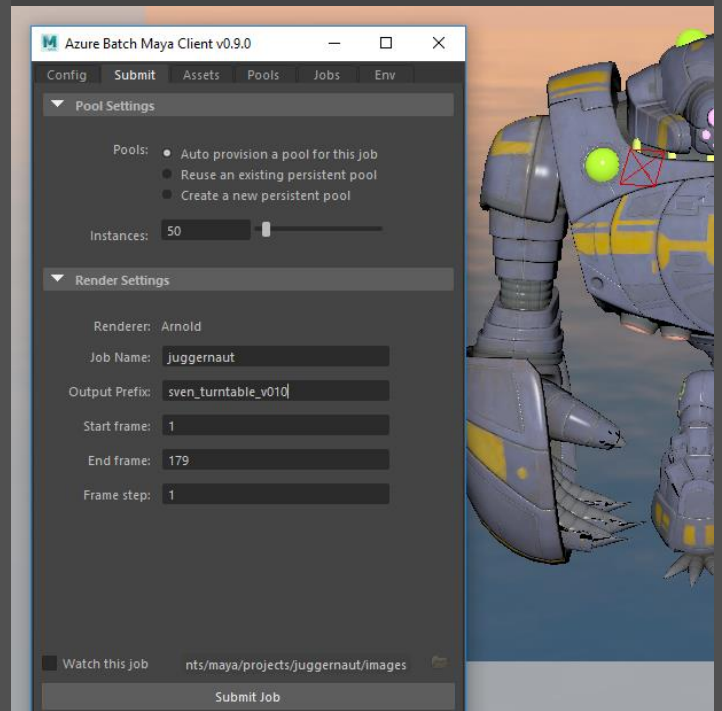

Azure Batch

AUTODESK® MAYA®

3DS MAX

SOLIDANGLE arnold

VM Renderer

VM Renderer

VM Renderer

Monitoring
Reporting
Single bill

Silicon alternatives

TRAINING
CPUs and GPUs, limited FPGAs,
ASICs under investigation

EVALUATION
CPUs and FPGAs,
ASICs under investigation

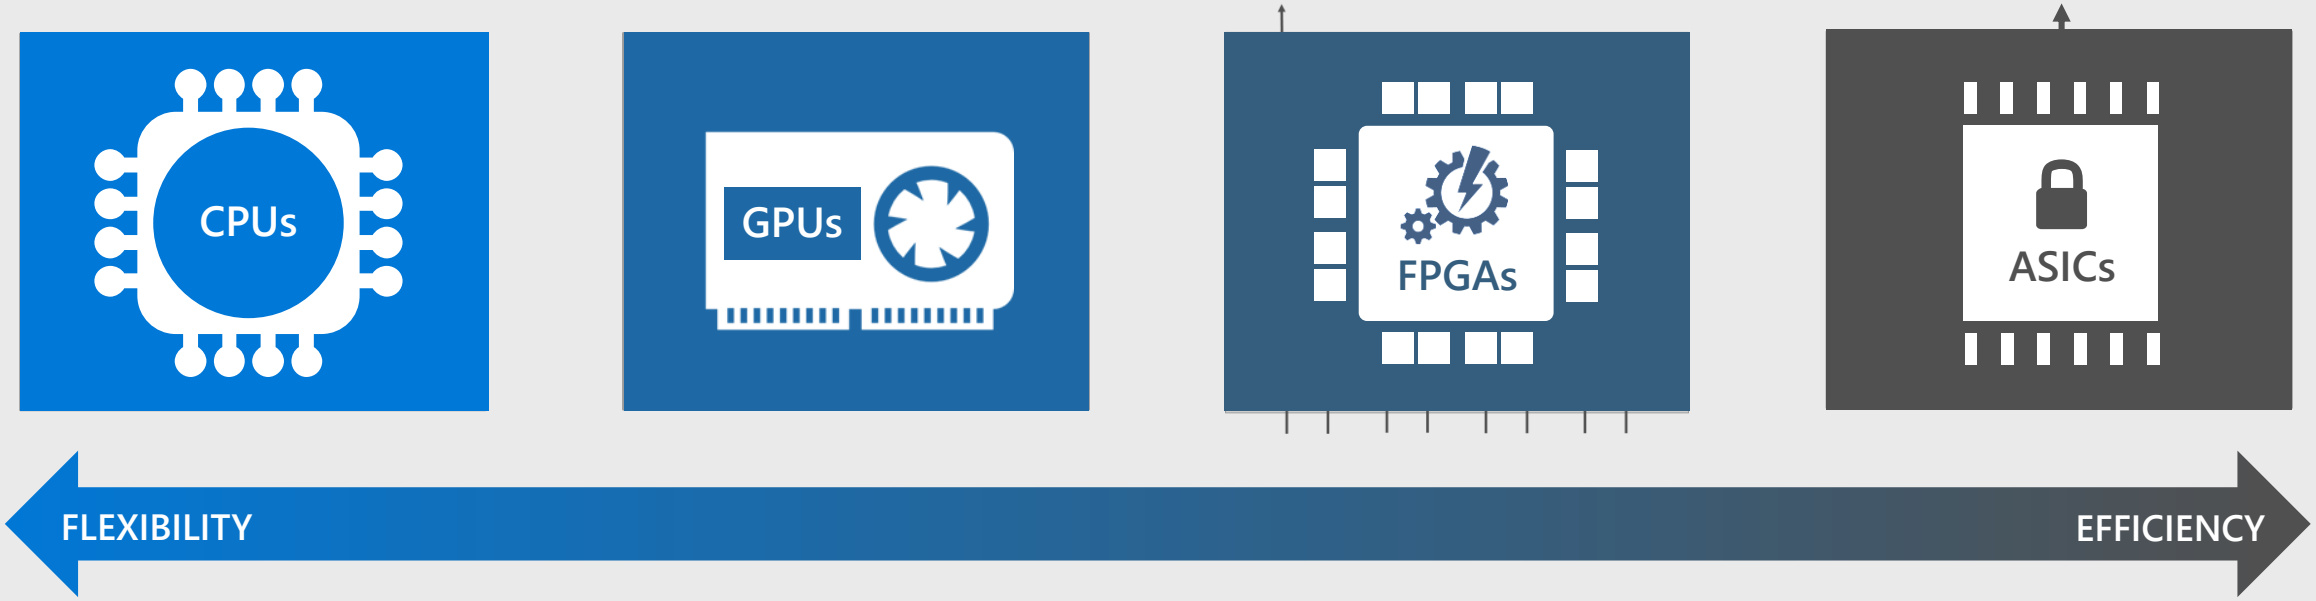

The power of deep learning on FPGA

Performance

- Tens to hundreds of TOPS of effective inference throughput at low batch sizes
- Ultra-low latency serving on modern DNNs
- >10X better than CPUs and GPUs
- Scale to many FPGAs in single DNN service

Flexibility

- FPGAs ideal for adapting to rapidly evolving ML
- CNNs, LSTMs, MLPs, reinforcement learning, feature extraction, decision trees, etc.
- Inference-optimized numerical precision
- Custom binarized, ternarized, tiny precision nets
- Sparsity, deep compression for larger, faster models

Scale

- Microsoft has the world's largest cloud investment in FPGAs
- Multiple Exa-Ops of aggregate AI capacity
- We have built powerful DNN serving platform on our FPGA fabric

Deep Learning Frameworks

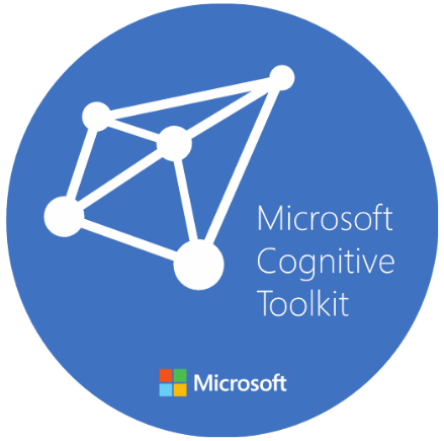

theano

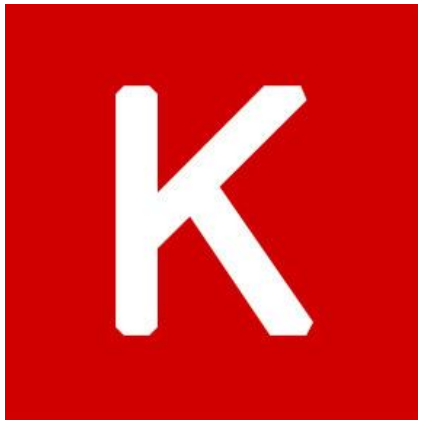

Caffe

Eu quero ensiar para o meu computador que isso é um gato

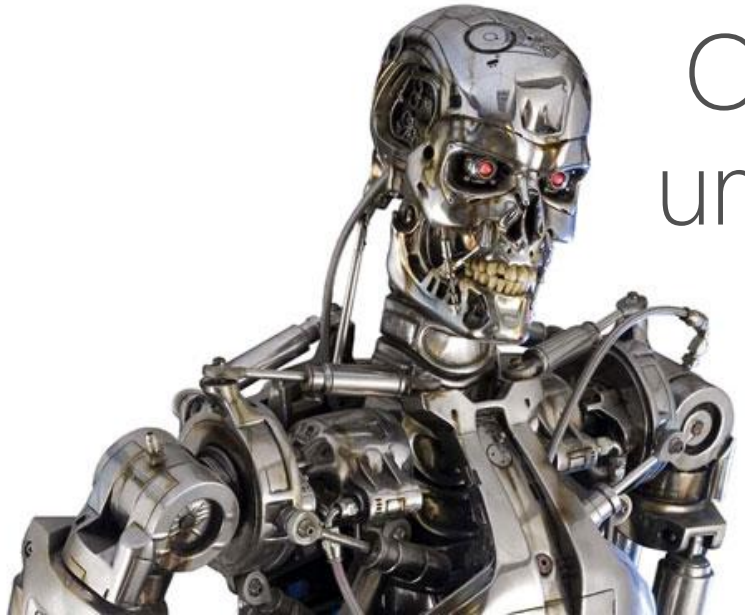

O que é
um gato?

Eu sou
um gato

Vamos ensinar esse computar

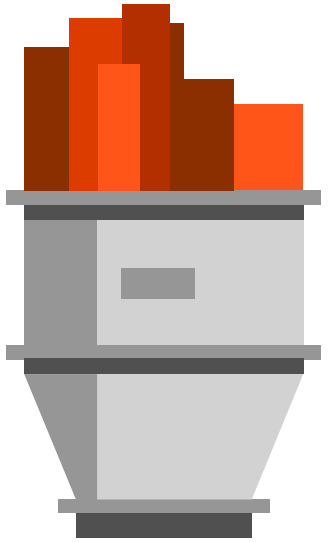

10101
01010
00100
11011
0101
010
00

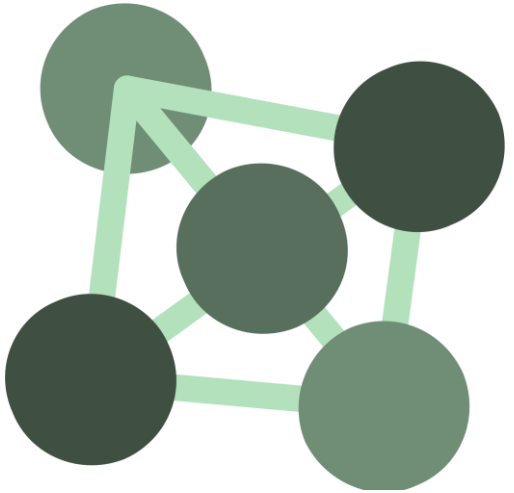

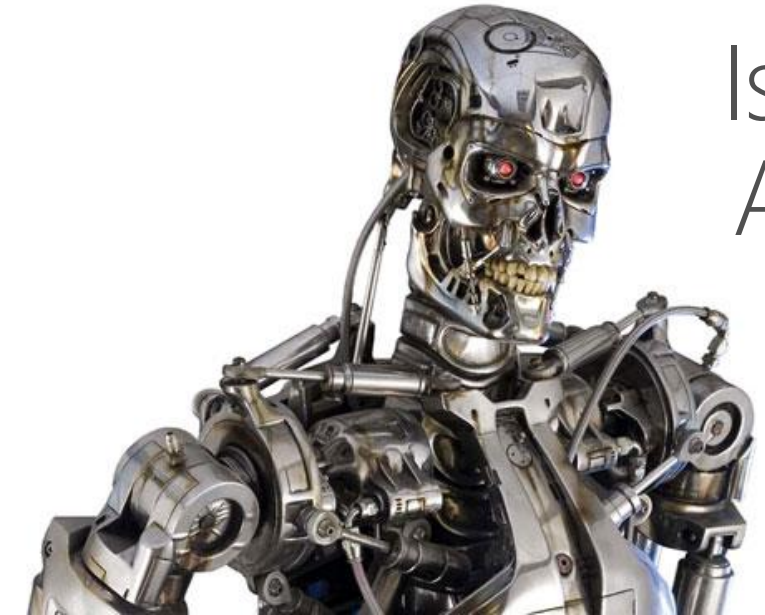

Isso é um gato.
Agora eu já sei.

WTF?!?!

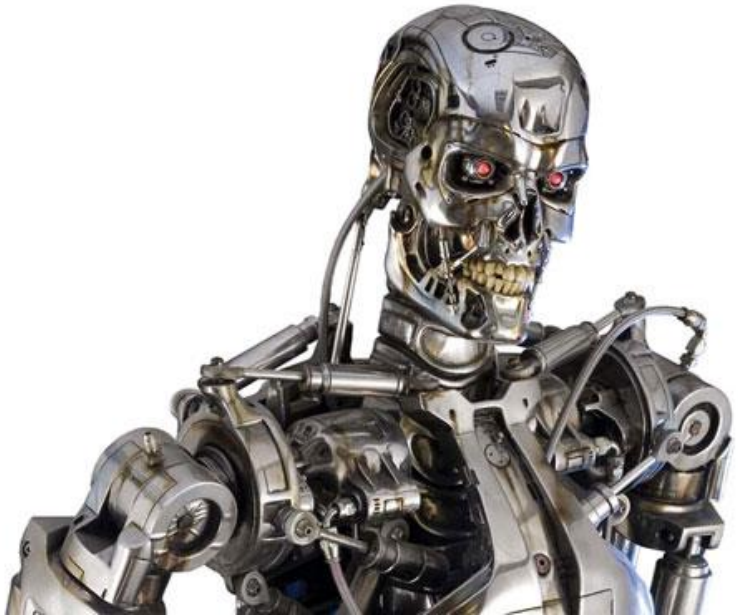

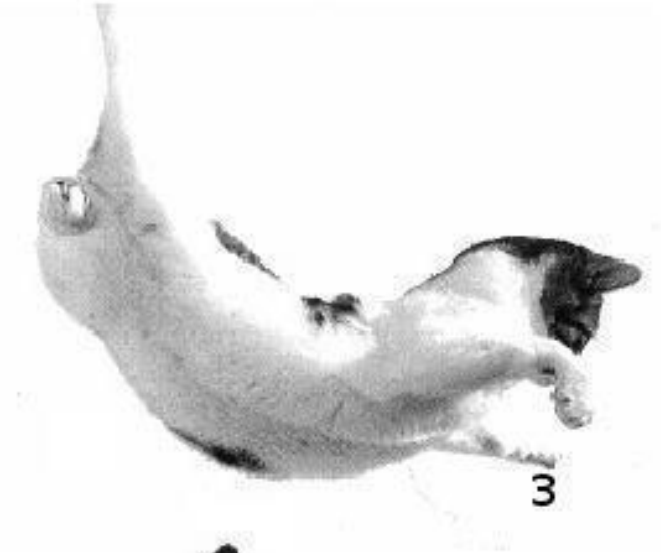

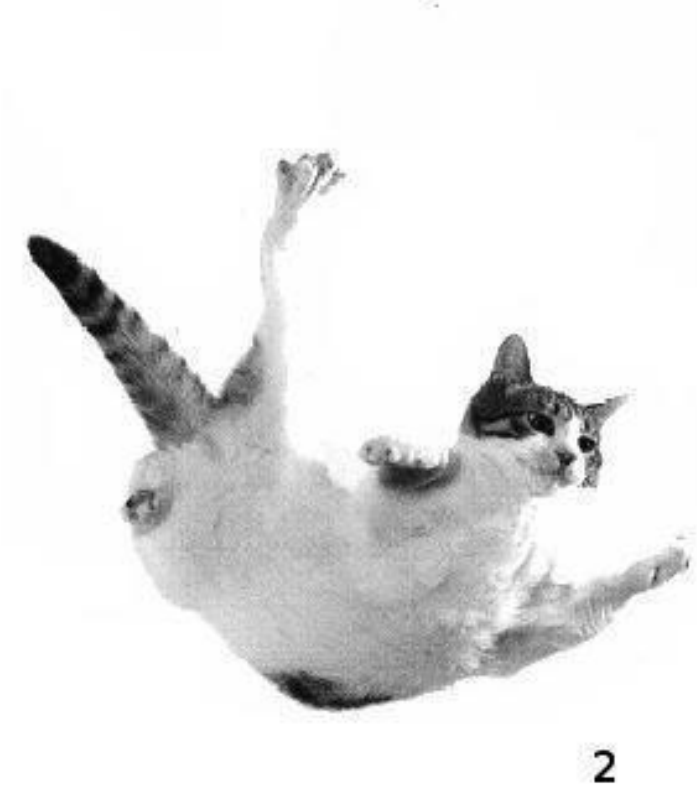

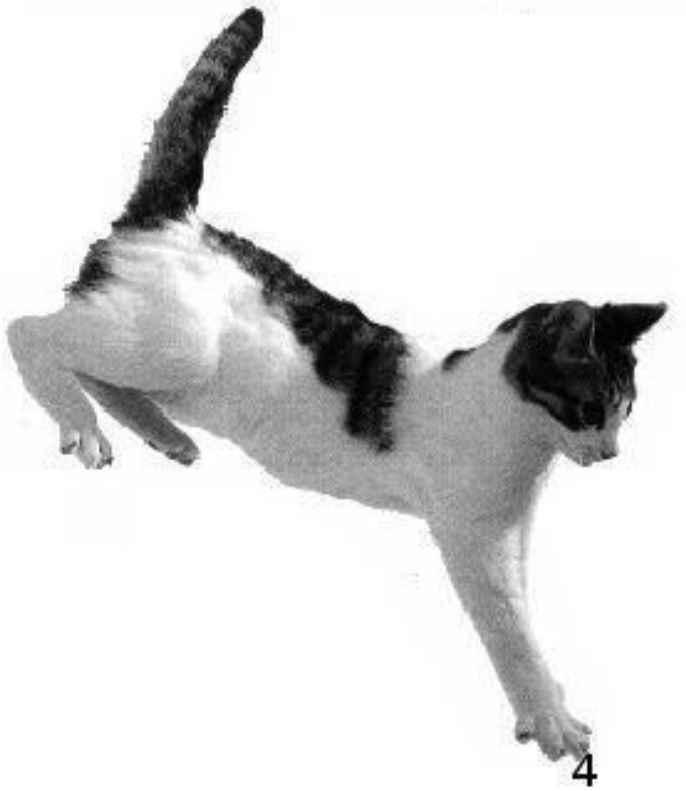

WTF?!?!

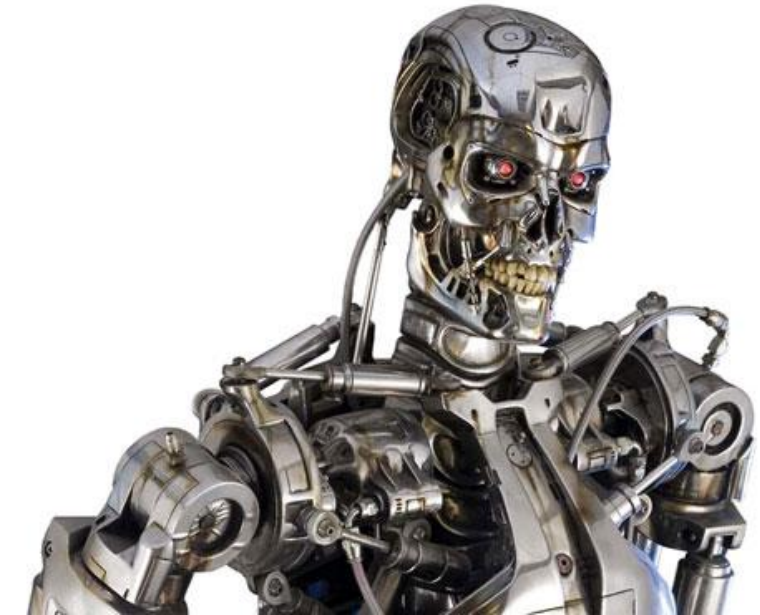

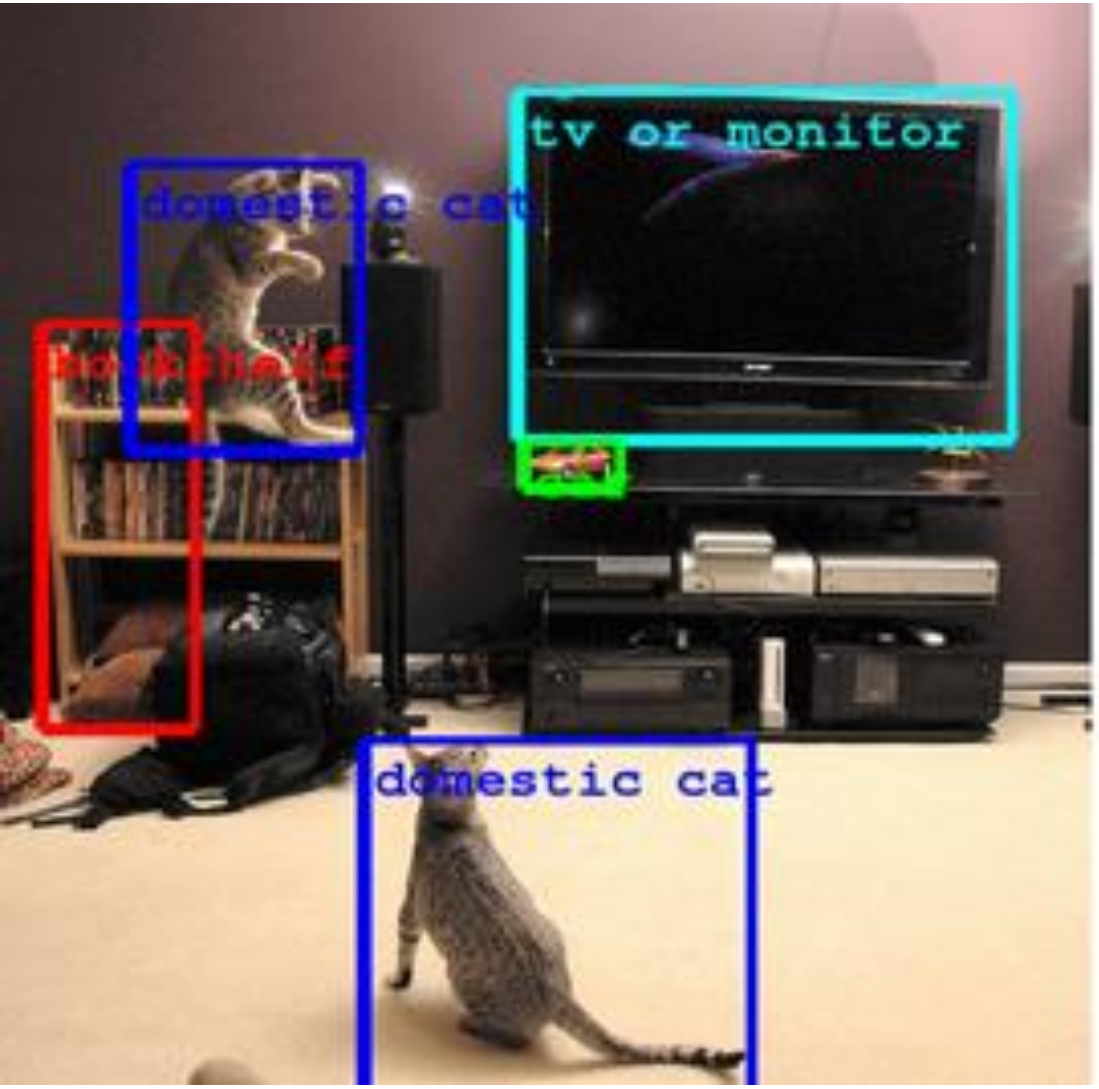
tv or monitor
domestic cat
domestic cat

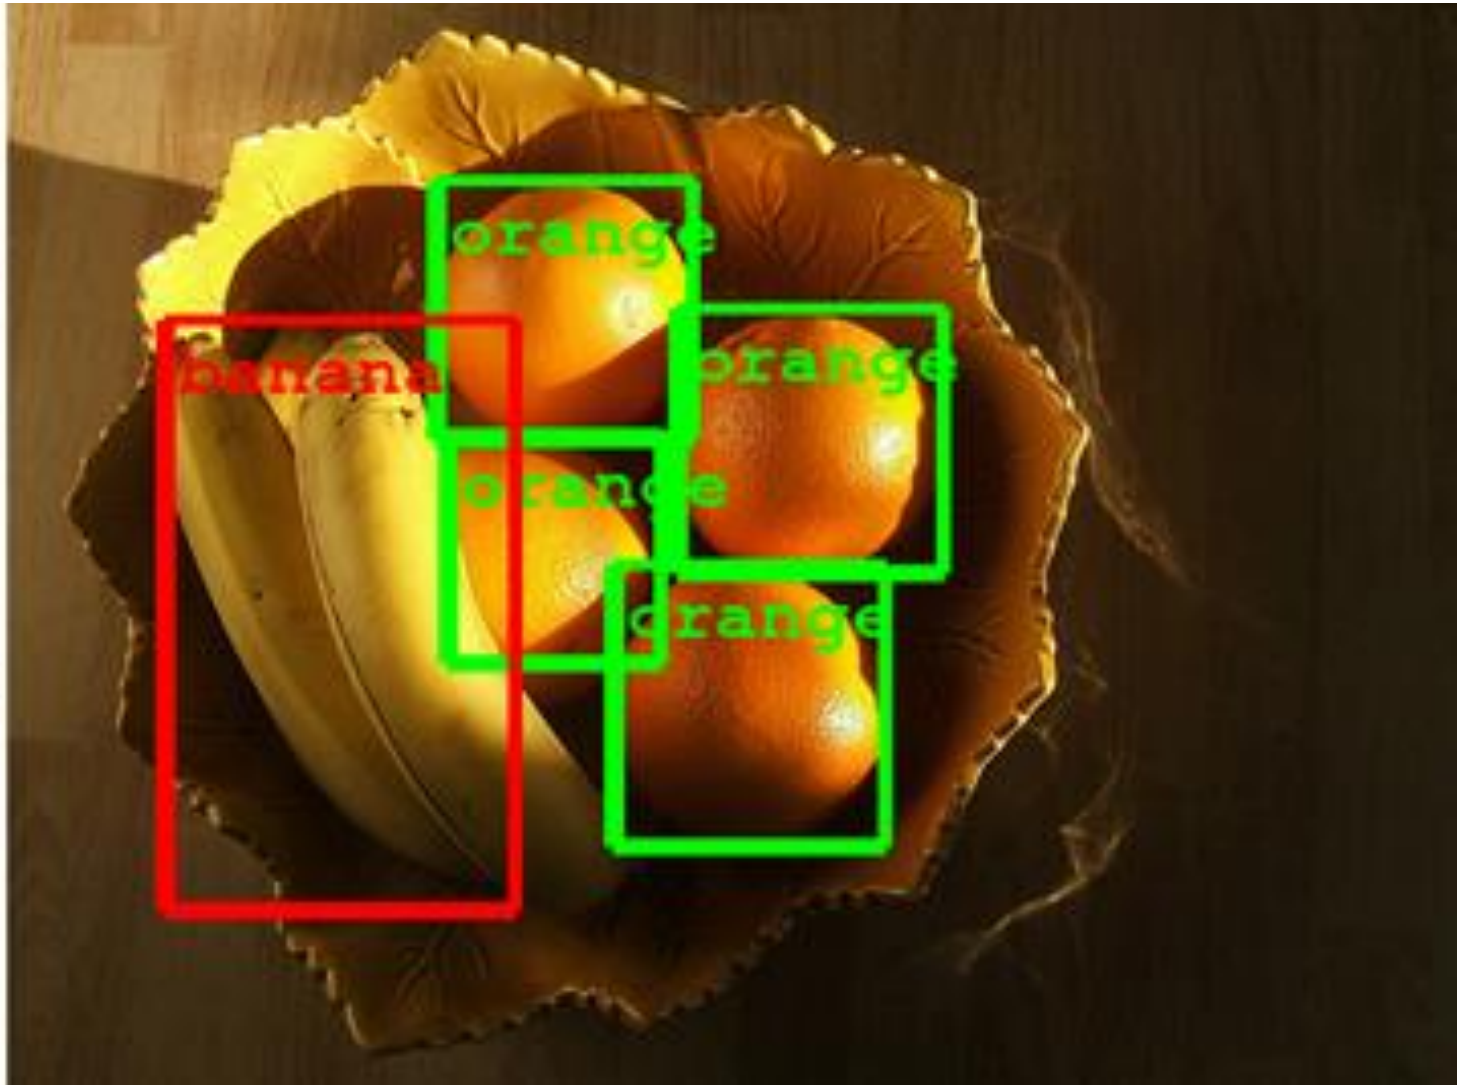
orange
orange
orange

Artificial Intelligence
Machine Learning
Deep Learning
Cognitive Computing

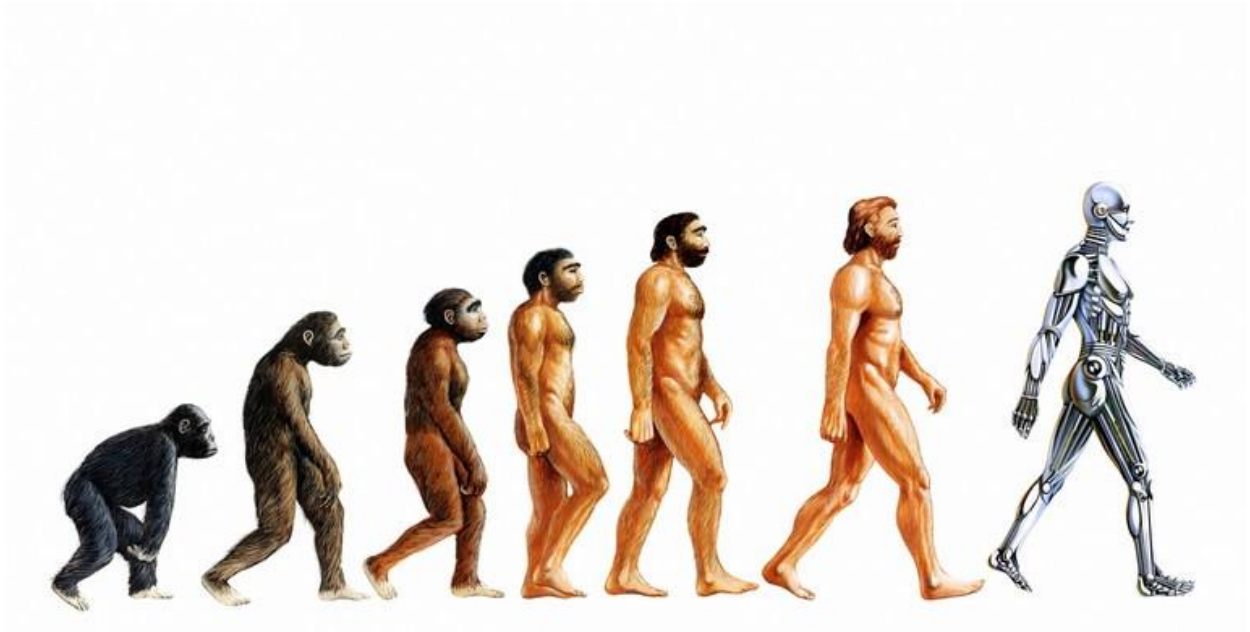

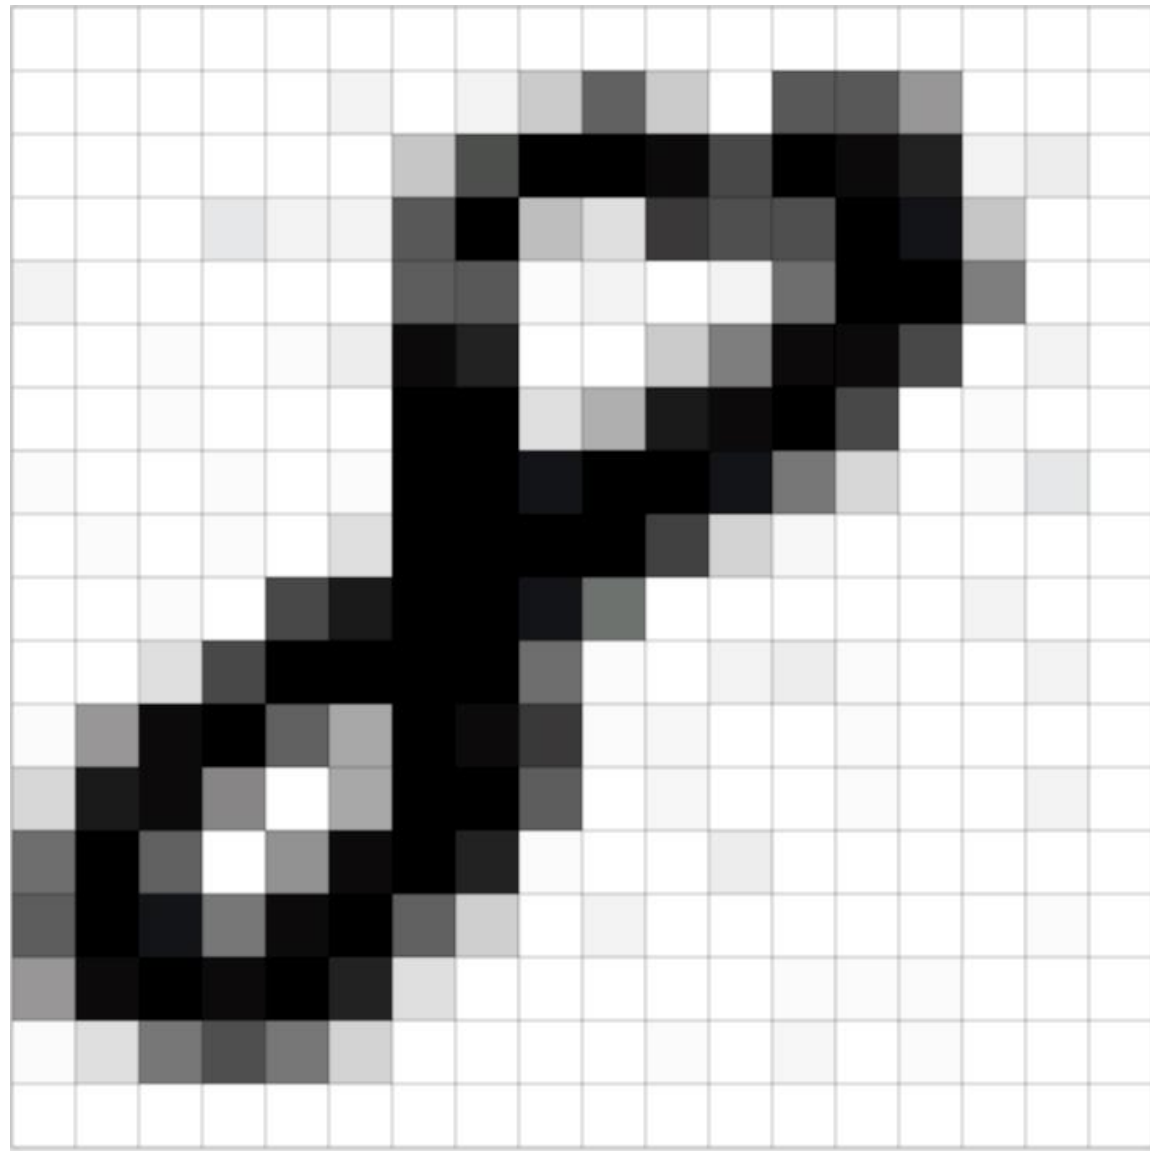

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	1	12	0	11	39	137	37	0	152	147	84	0	0	0
0	0	1	0	0	0	41	160	250	255	235	162	255	238	206	11	13	0
0	0	0	16	9	9	150	251	45	21	184	159	154	255	233	40	0	0
10	0	0	0	0	0	145	146	3	10	0	11	124	253	255	107	0	0
0	0	3	0	4	15	236	216	0	0	38	109	247	240	169	0	11	0
1	0	2	0	0	0	253	253	23	62	224	241	255	164	0	5	0	0
6	0	0	4	0	3	252	250	228	255	255	234	112	28	0	2	17	0
0	2	1	4	0	21	255	253	251	255	172	31	8	0	1	0	0	0
0	0	4	0	163	225	251	255	229	120	0	0	0	0	0	11	0	0
0	0	21	162	255	255	254	255	126	6	0	10	14	6	0	0	9	0
3	79	242	255	141	66	255	245	189	7	8	0	0	5	0	0	0	0
26	221	237	98	0	67	251	255	144	0	8	0	0	7	0	0	11	0
125	255	141	0	87	244	255	208	3	0	0	13	0	1	0	1	0	0
145	248	228	116	235	255	141	34	0	11	0	1	0	0	0	1	3	0
85	237	253	246	255	210	21	1	0	1	0	0	6	2	4	0	0	0
6	23	112	157	114	32	0	0	0	0	2	0	8	0	7	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

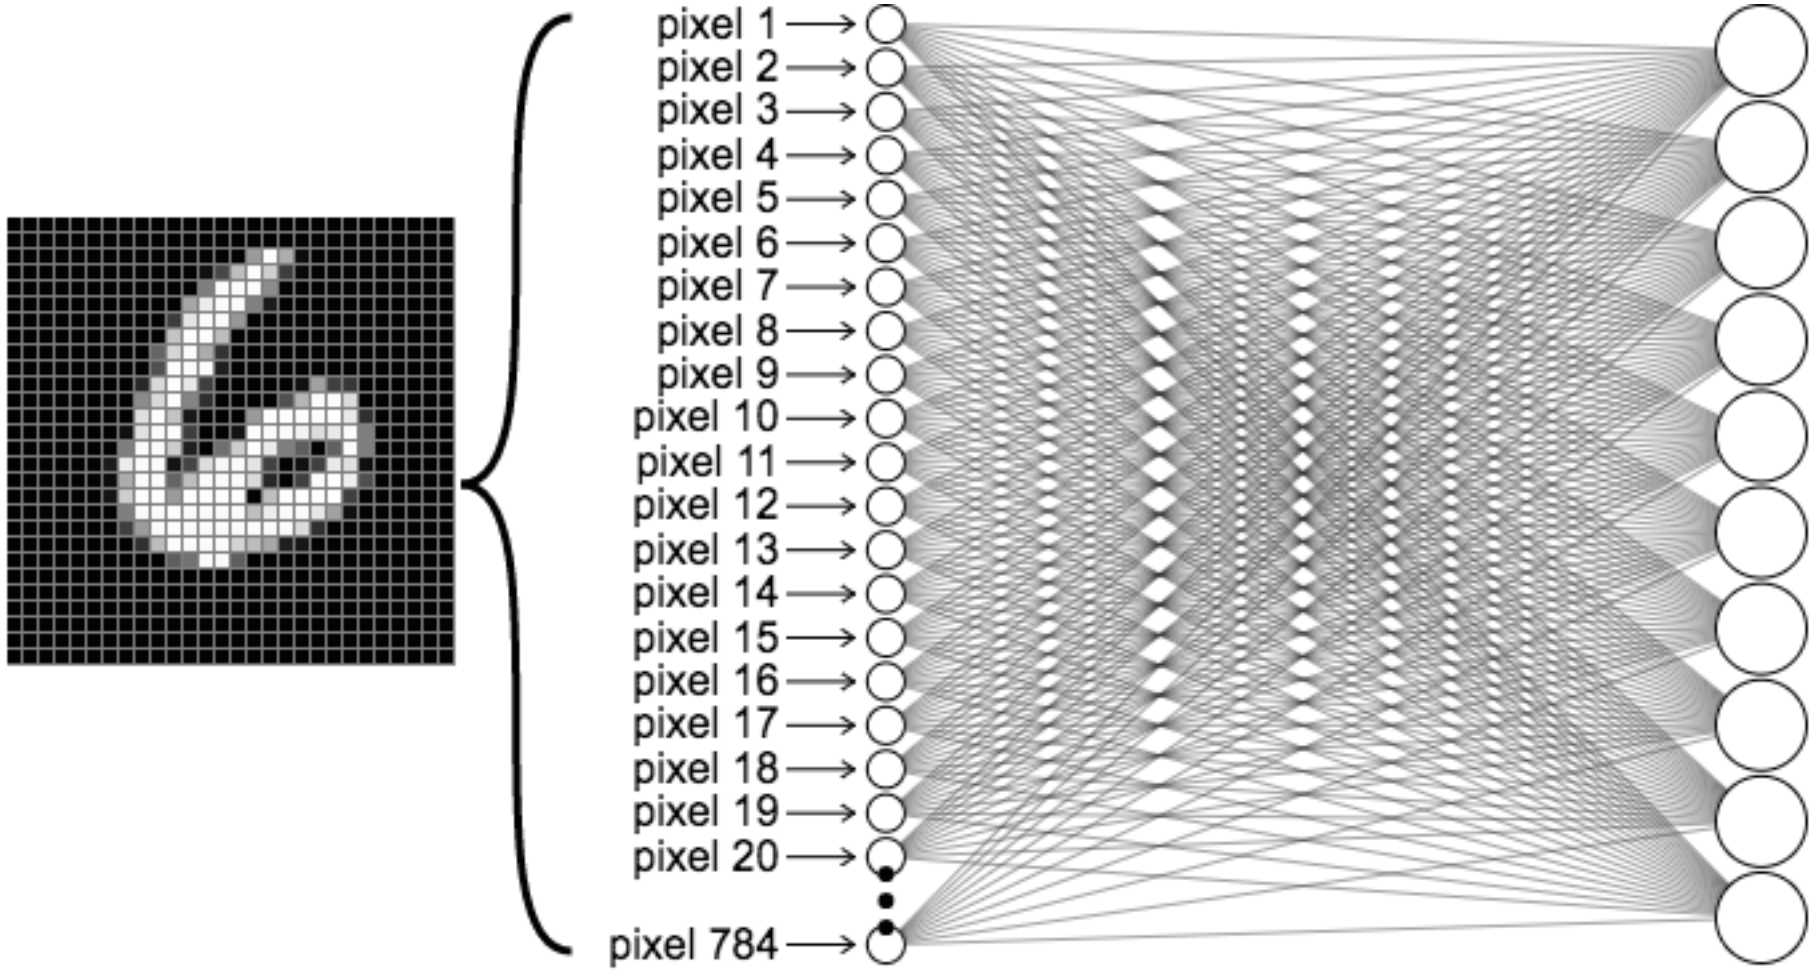
pixel 1
pixel 2
pixel 3
pixel 4
pixel 5
pixel 6
pixel 7
pixel 8
pixel 9
pixel 10
pixel 11
pixel 12
pixel 13
pixel 14
pixel 15
pixel 16
pixel 17
pixel 18
pixel 19
pixel 20
pixel 784

Escalar

s = np.array(8)

s.shape = ( )

[8]

Vetor

v = np.array([1,2,3])

v.shape = (3,)

[1,2,3]

Matriz

m = np.array([[1,2,3], [4,5,6], [7,8,9]])

m.shape = (3, 3)

[1,2,3]

[4,5,6]

[7,8,9]

Tensor

```
t = np.array([[[[1],[2]],[[3],[4]],[[5],[6]]],
              [[[7],[8]],[[9],[10]],[[11],[12]]],
              [[[13],[14]],[[15],[16]],[[17],[17]]]])
```

t.shape = (3, 3, 2, 1)

't'
'e'
'n'
's'
'o'
'r'

Vetor

3	1	4	1
5	9	2	6
5	3	5	8
9	7	9	3
2	3	8	4
6	2	6	4

Matriz

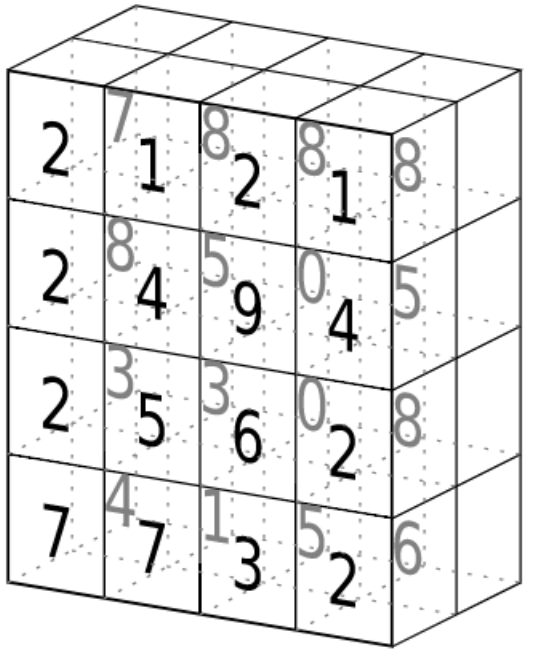

Tensor

Perceptron

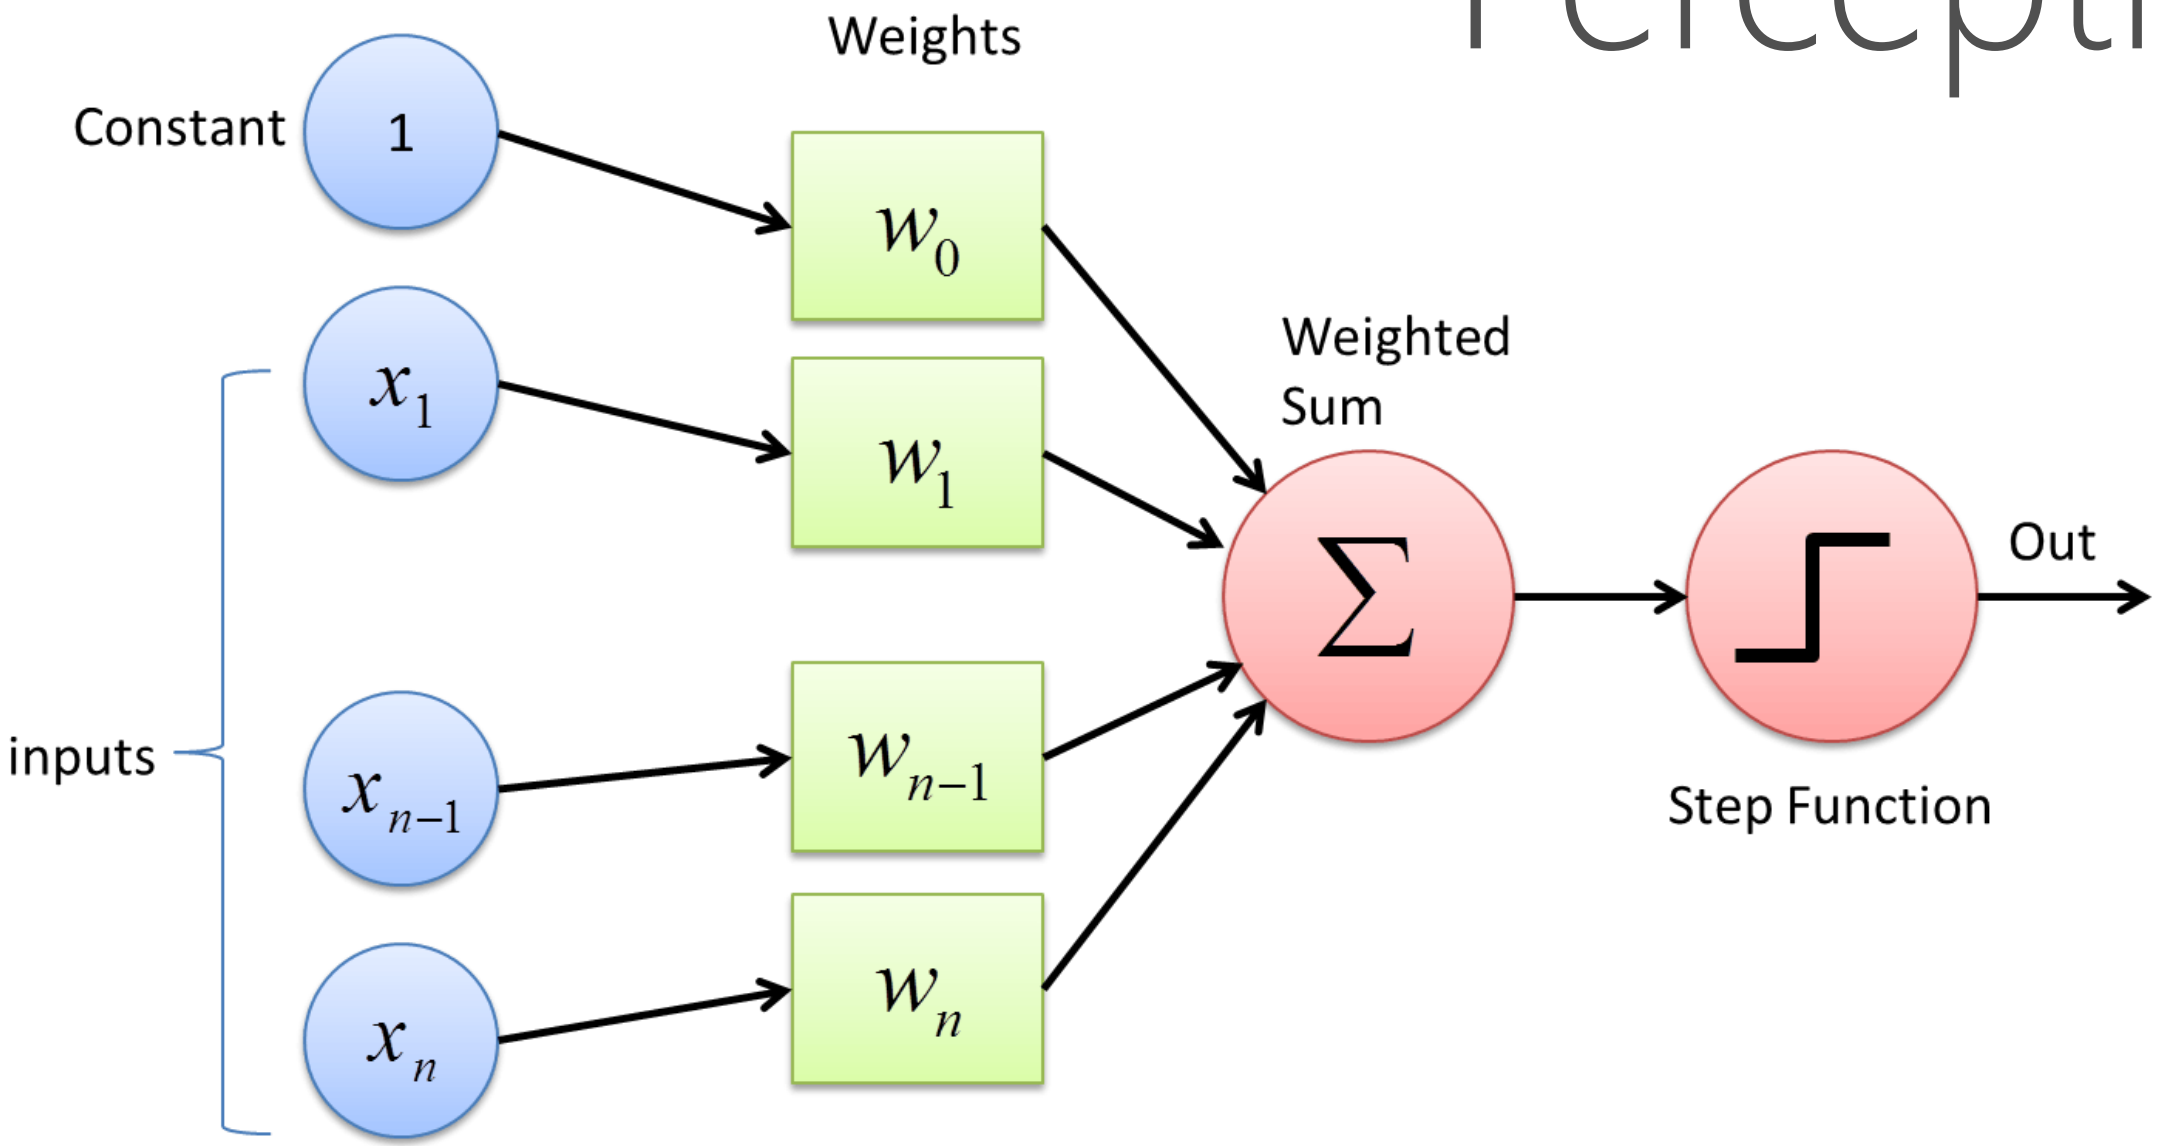
Weights
Constant
1
w_0
x_1
w_1
Weighted
Sum
Σ
Out
inputs
x_{n-1}
w_{n-1}
x_n
w_n
Step Function

What?

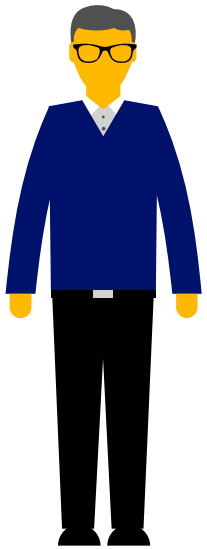

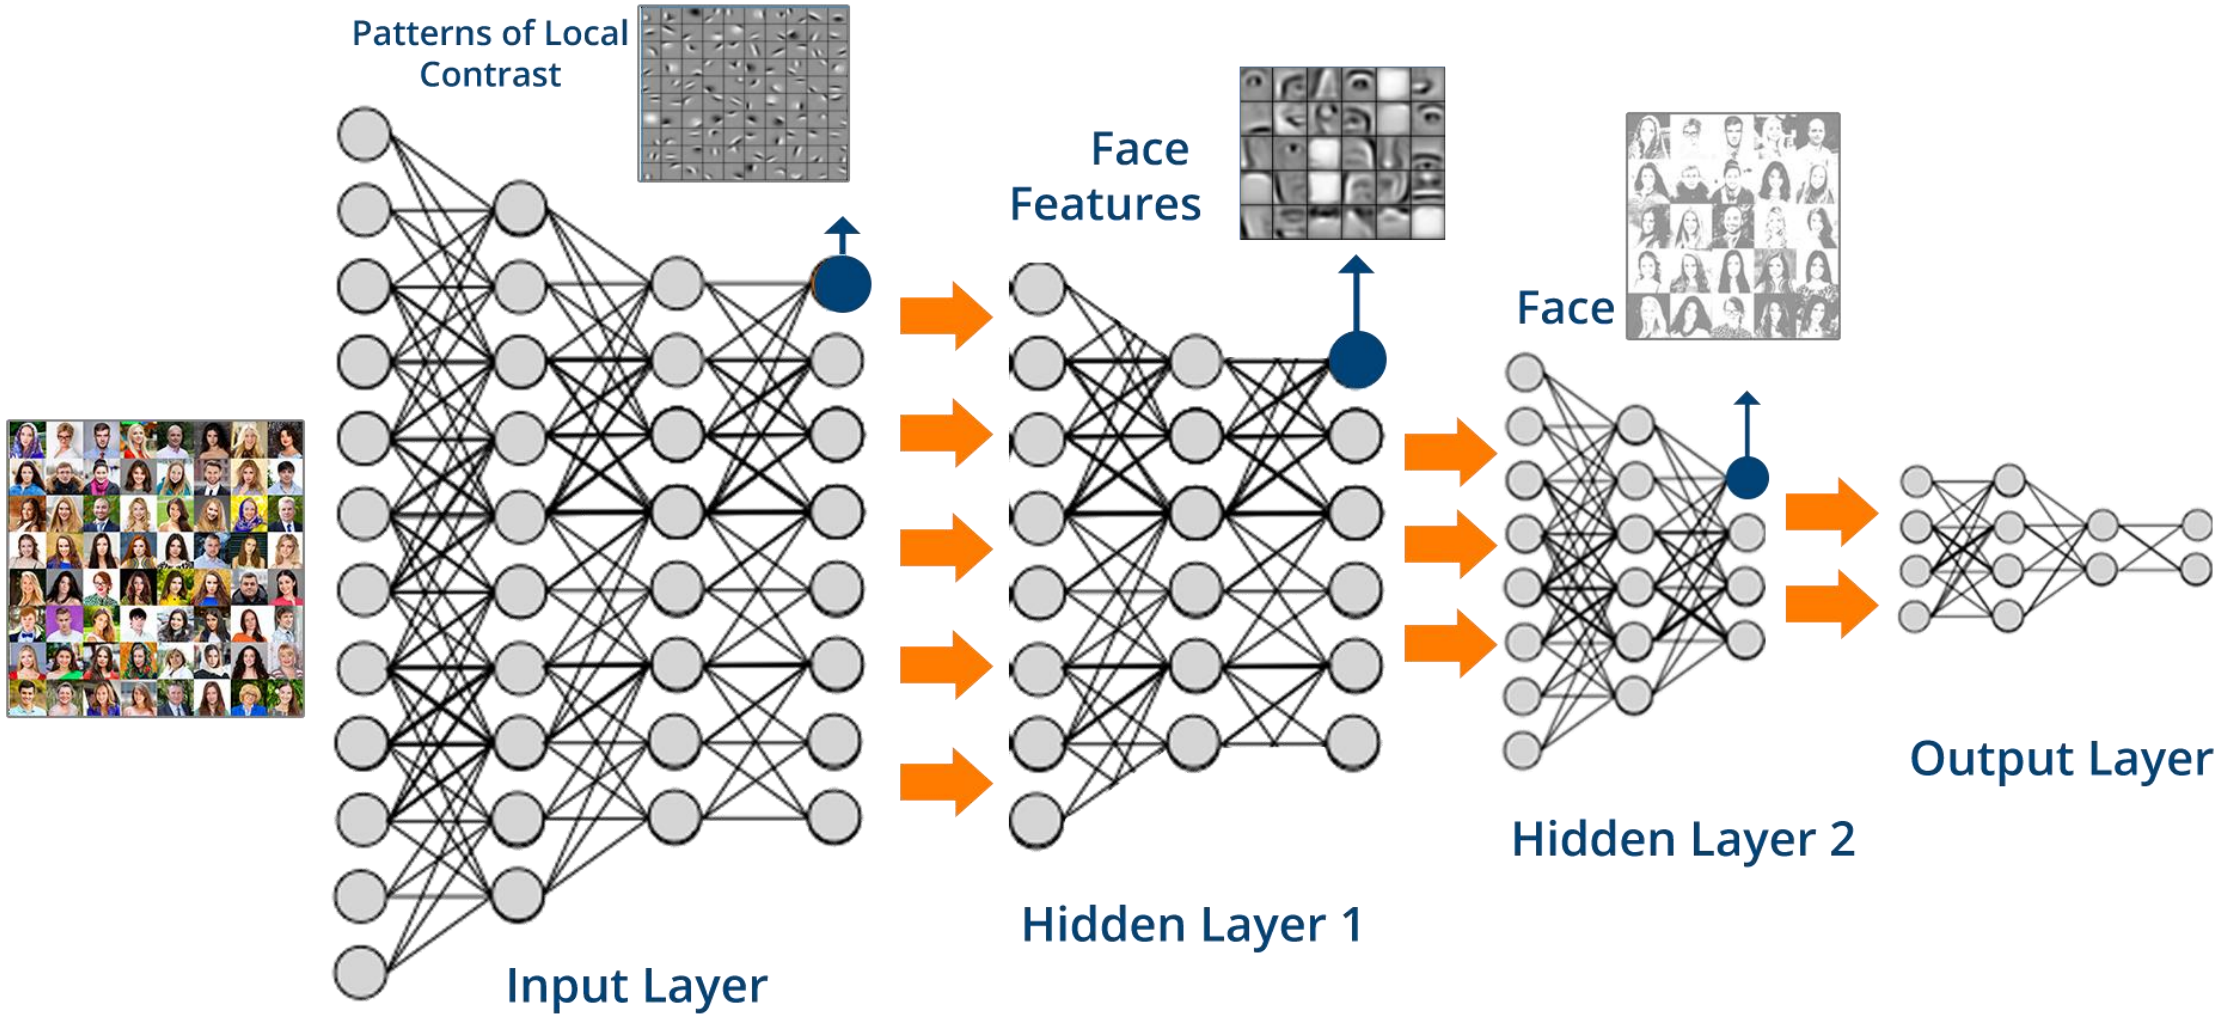
Patterns of Local Contrast
Face Features
Face
Input Layer
Hidden Layer 1
Hidden Layer 2
Output Layer

Generative Adversarial Networks

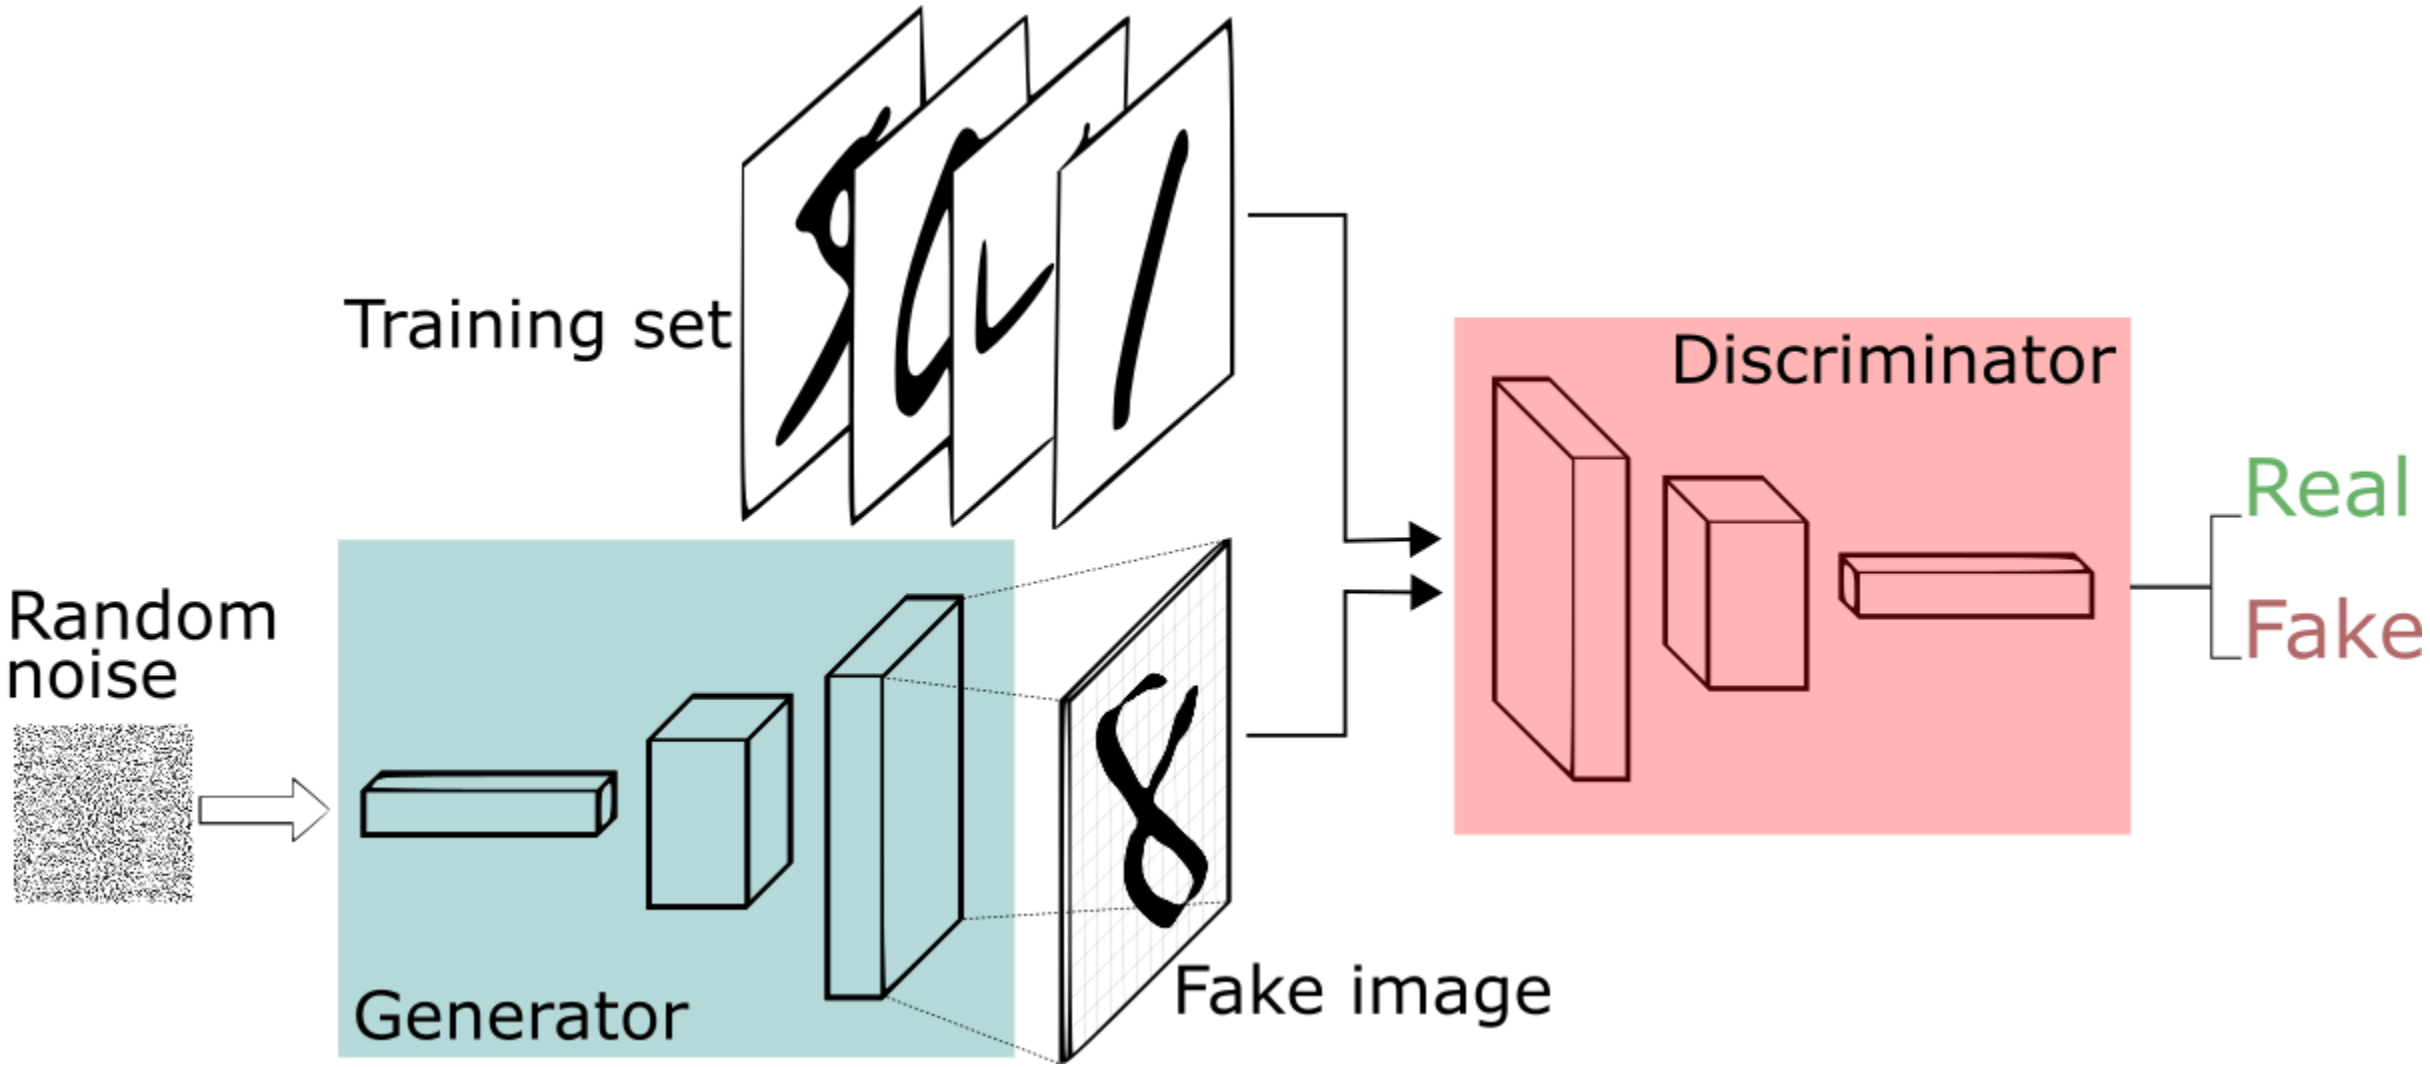

Multiple GPU Cards

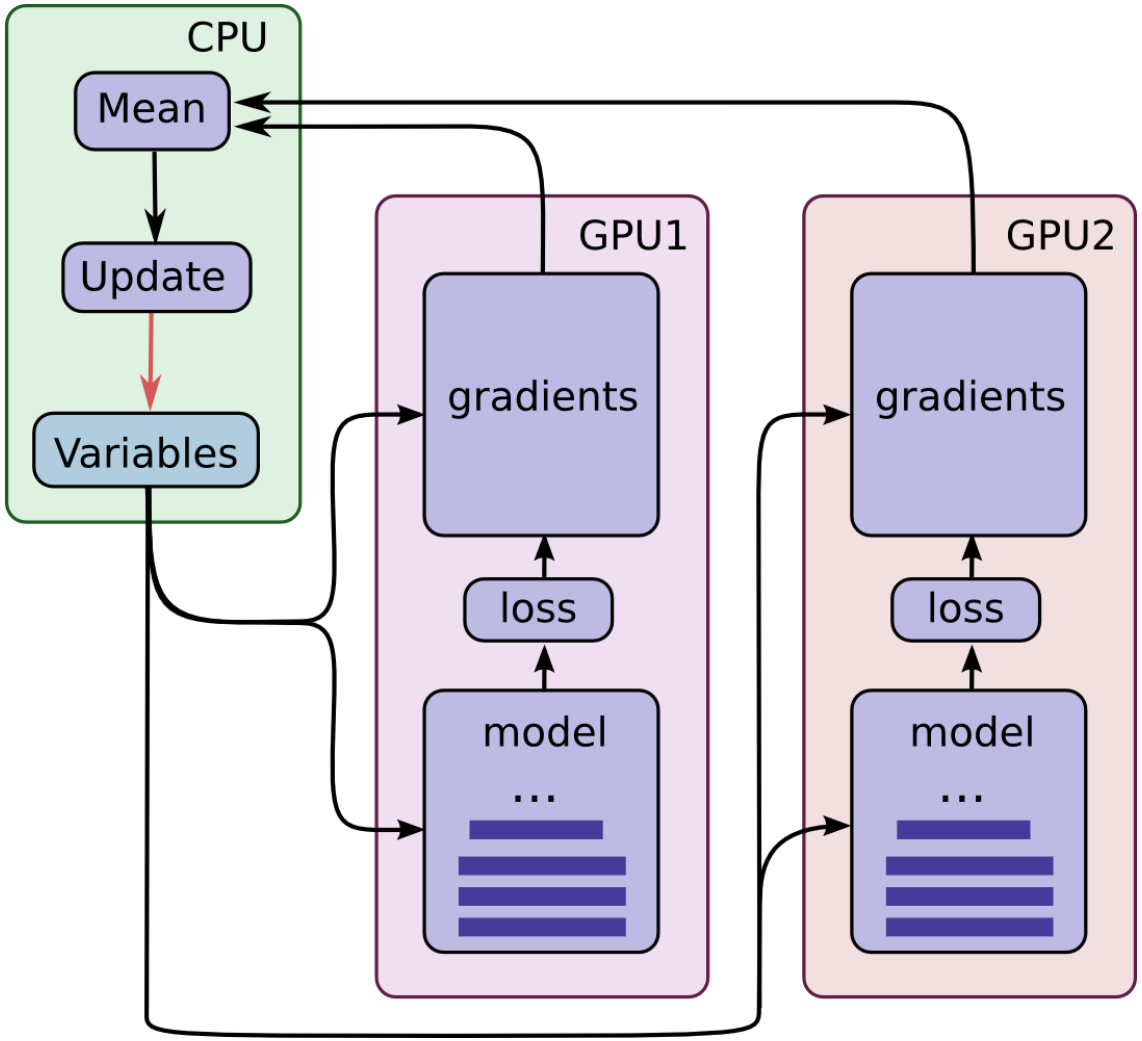

Parallel Training

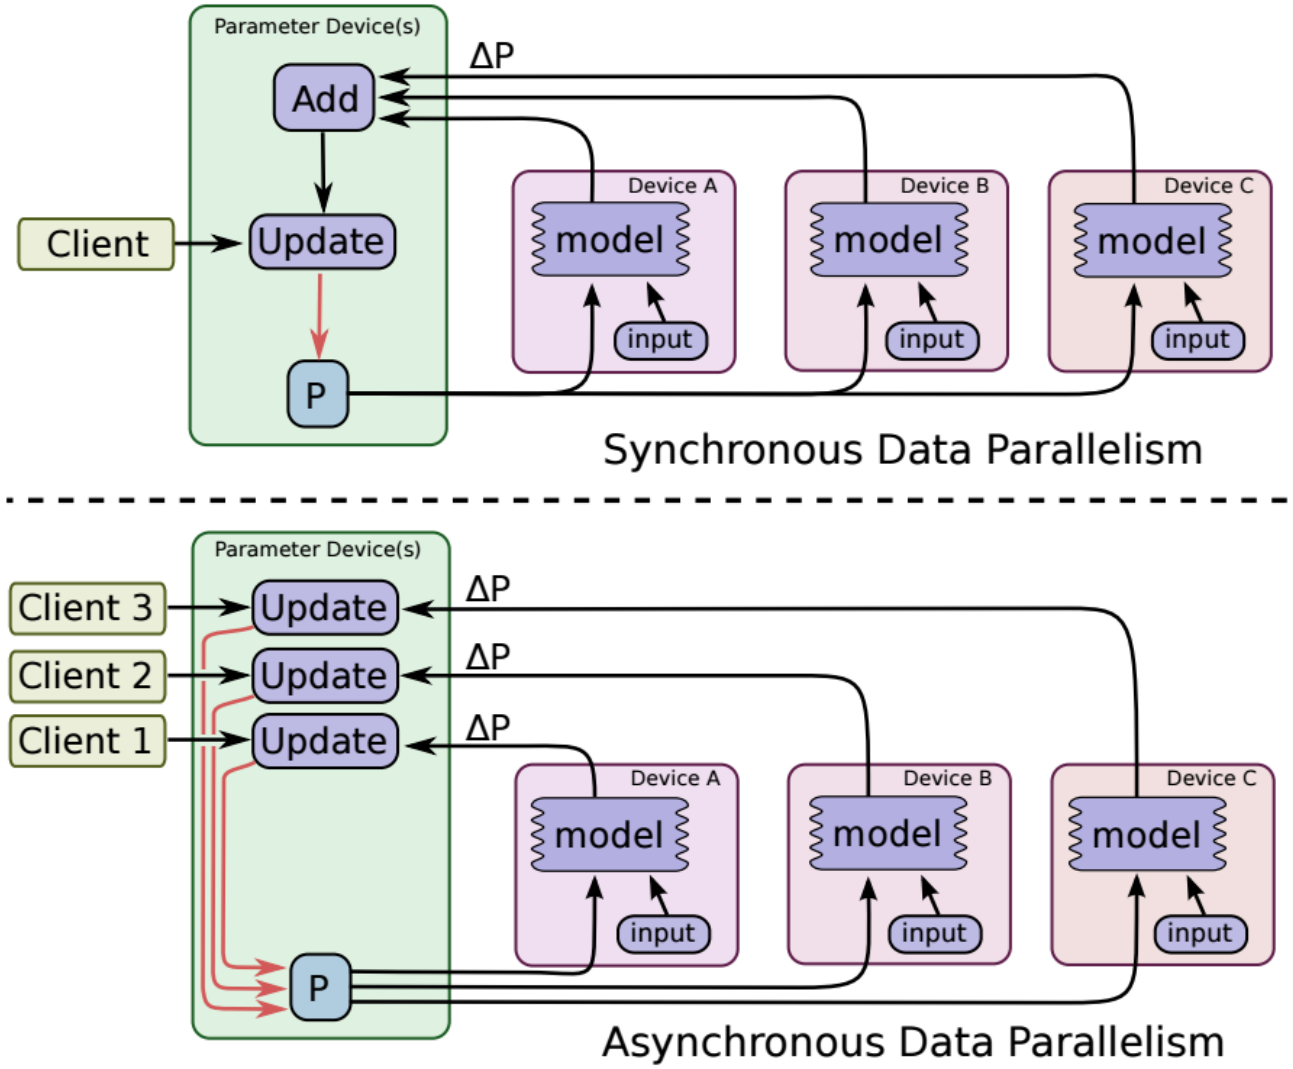

1 SHOW
2 ME
3 THE
4 CODE

Thank you!

2018

Global Azure

BOOTCAMP

vitormeriat.com.br

linkedin.com/in/vitormeriat

twitter.com/vitormeriat

github.com/vitormeriat

youtube.com/vitormeriat